

It is my pleasure to recommend Olga Pearce for the Professor of Practice position at Texas A&M University.

. . . about the recommender here ...

I have known Olga through the conference committee work for nearly a decade, and have worked closely with her over the past year on the Hatchet project.

Olga is the PI for the Hatchet project, a performance analysis tool that enables users to programmatically compare multiple performance profiles, and write their own analysis. Hatchet is a young two year old project, but has already attracted many users. As a tools researcher and a performance analysis expert working with many large scale parallel simulation codes, Olga was instrumental in developing the initial idea for Hatchet to fill the need for programmatic performance analysis.

While it was important to develop a tool that would work with existing profiling and tracing tools, Olga early on recognized the importance of interfacing with the in-house profiling tool called Caliper, which Olga has been helping integrate into several of LLNL's simulation codes. Olga was instrumental in integrating Hatchet into Spot, a web interface code developers use to view the performance data Caliper collects about executions of their simulations. Because code developers rely on Olga as the tools expert, they would quickly call her to find out the answer to "what does this button do?", at which point she would recruit them as early users.

As feature requests poured in, Olga recruited several students to help the otherwise small Hatchet team. At this point, Hatchet has two LLNL staff (counting Olga) working on it, three university professors (from different institutions) with one or two students each. The team now includes experts in simulation optimization, performance tools, and visualization, enabling creative solutions that are rapidly delivered to the very engaged users of the tool.

While the early adoption of Hatchet was by word of mouth, Olga has since lead several Hatchet tutorials at LLNL. Olga proactively talks to the current and prospective Hatchet users, bringing back feature requests and ideas to the Hatchet team. Olga facilitates priorities for the project, both near and long term, as well as software releases and coordination with sister-projects. Additionally, Olga helps mentor the students on the project, along with the junior staff member.

In the wider Computer Science community, Olga is very active in working with students, over the years serving on the Cluster Challenge, Broader Engagement, and Student committees at the ACM/IEEE Supercomputing Conference, and the Ph.D. forum at the International Conference on Su-

percomputing. Olga's technical expertise is widely recognized, and Olga has been invited to and serves on prominent technical committees in her field, such as the ACM/IEEE Supercomputing Conference, IEEE International Parallel & Distributed Processing Symposium, EuroMPI, IEEE International Conference on Cluster Computing, International Conference on Parallel Processing, IEEE/ACM International Conference on Cluster, Cloud, & Grid Computing, International Conference on Supercomputing, as well as a reviewer for several journals. Additionally, Olga co-chairs the Salishan Conference on High Speed Computing, guiding the direction of this premier invitation-only conference on computer architecture, languages, and algorithms organized by the tri-labs (LLNL, LANL, Sandia).